

Gigahorse: Thorough, Declarative Decompilation of Smart Contracts

Neville Grech
University of Athens
and University of Malta
Greece and Malta
me@nevillegrech.com

Lexi Brent
The University of Sydney
Australia
lexi.brent@sydney.edu.au

Bernhard Scholz
The University of Sydney
Australia
bernhard.scholz@sydney.edu.au

Yannis Smaragdakis
University of Athens
Greece
yannis@smaragd.org

Abstract—The rise of smart contracts—autonomous applications running on blockchains—has led to a growing number of threats, necessitating sophisticated program analysis. However, smart contracts, which transact valuable tokens and cryptocurrencies, are compiled to very low-level bytecode. This bytecode is the ultimate semantics and means of enforcement of the contract.

We present the Gigahorse toolchain. At its core is a reverse compiler (i.e., a decompiler) that decompiles smart contracts from Ethereum Virtual Machine (EVM) bytecode into a high-level 3-address code representation. The new intermediate representation of smart contracts makes implicit data- and control-flow dependencies of the EVM bytecode explicit. Decompilation obviates the need for a contract’s source and allows the analysis of both new and deployed contracts.

Gigahorse advances the state of the art on several fronts. It gives the highest analysis precision and completeness among decompilers for Ethereum smart contracts—e.g., Gigahorse can decompile over 99.98% of deployed contracts, compared to 88% for the recently-published Vandal decompiler and under 50% for the state-of-the-practice Porosity decompiler. Importantly, Gigahorse offers a full-featured toolchain for further analyses (and a “batteries included” approach, with multiple clients already implemented), together with the highest performance and scalability. Key to these improvements is Gigahorse’s use of a declarative, logic-based specification, which allows high-level insights to inform low-level decompilation.

Index Terms—Ethereum, Blockchain, Decompilation, Program Analysis, Security

I. INTRODUCTION

Distributed blockchain platforms have captured the imagination of scientists and the public alike. Blockchain technology offers decentralized consensus mechanisms for any transactions that, in the past, would have required a trusted centralized authority. One of the most evident embodiments of this vision is the development of *smart contracts*: Turing-complete autonomous agents that run on distributed blockchains, such as Ethereum or Cardano. A smart contract may, for instance, implement a lending policy, a charging scheme for digital goods, an auction, the full set of operations of a bank, and virtually any other logic governing multi-party transactions.

Ethereum is the best-known, most popular blockchain platform that supports full-featured smart contracts. (As of this writing, the Ethereum cryptocurrency market capitalization is \$13B.) Ethereum offers an excellent demonstration of the potential for smart contracts, as well as their technical challenges. Developers typically write smart contracts in a high-level language called Solidity, which is compiled into

immutable low-level Ethereum VM (EVM) bytecode for the blockchain’s distributed virtual machine.

The open nature of smart contracts, as well as their role in handling high-value currency, raise the need for thorough contract analysis and validation. This task is hindered, however, by the low-level stack-based design of the EVM bytecode that has hardly any abstractions as found as in other languages, such as Java’s virtual machine. For example, there is no notion of functions or calls—a compiler that translates to EVM bytecode needs to invent its own conventions for implementing local calls over the stack.

It is telling that recent research [1], [9], [22], [34] has focused on decompiling smart contracts into a higher-level representation, before applying any further (usually security-oriented) analysis. Past decompilation efforts have been, at best, incomplete. The best-known decompiler (largely defining the state-of-the-practice) is Porosity [33], which in our study fails to yield results for 50% of deployed contracts of all smart contracts on the block chain. Upcoming research tools including the Vandal decompiler [35] still fail to decompile a significant portion of real contracts (around 12%) due to the complex task of converting EVM’s stack-based operations to a register-based intermediate representation.

Such difficulties are much more than technicalities of the platform or idiosyncrasies of existing tools. Any current or future smart contract platform is likely to employ virtual machines that are low-level. The designs of these virtual machines are optimized for massively replicated execution of smart contracts. The bytecode effectively represents an assembly language designed for efficient execution and compact program representation, since the bytecode must be stored on the blockchain. Hence, the bytecode for smart contracts will never be optimal for human readability or reverse compilation.

For effective decompilation, significant program comprehension of the bytecode is necessary. A decompiler requires deep program *understanding* before it can reconstruct the bytecode to a high-level representation. For instance, to recognize which low-level jumps correspond to high-level function calls, a decompiler must deduce possible addresses for jump instructions to be able to reconstruct the control-flow of a smart contract. Such understanding is compounding: once calls are recognized as calls (and not just as a mere intra-procedural jump), its decompilation precision can further improve, by pruning impossible targets.

In this paper, we introduce the Gigahorse toolchain for analysis of Ethereum smart contracts. Gigahorse addresses the above challenges, making the following contributions:

- It offers a highly-effective decompiler, yielding higher precision and completeness than the state-of-the-art (e.g., decompiling virtually all existing contracts on the Ethereum blockchain).
- Gigahorse anchors around it a full-featured tool suite, offering libraries for building analyses, as well as ready-made clients for existing analyses.
- The Gigahorse approach provides new decompilation insights (possibly of value for many more platforms): As higher level program features are discovered, they feed back to lower level analyses. E.g., by discovering functions, stack analyses are performed locally (more precisely), and the effects of function calls on the stack are also summarized, enabling both more precise and more scalable analysis.
- Gigahorse showcases an unconventional decompilation approach, which enables the above benefits: the decompiler is specified declaratively, using logic-based (Datalog) rules.
- Gigahorse is evaluated and its features illustrated on the full set of smart contracts of the Ethereum blockchain.

Gigahorse is powering the ongoing free Contract Library service (<https://contract-library.com>), which offers decompiled versions of all contracts on the Ethereum blockchain. Contract Library is a valuable service for Ethereum security analysts, and is currently receiving several tens of unique visitors and over a thousand page views per day.

II. BACKGROUND

We next present some background on the EVM bytecode language and declarative static analysis.

A. Low-Level Bytecode in the Ethereum Blockchain

The EVM is a stack-based low-level intermediate representation (IR). In the bytecode form of a smart contract, symbolic information has been replaced by numeric constants, functions are fused together in a sea of instructions, and control-flow is obfuscated by jump addresses that are popped from the stack. To highlight this issue, it is instructive to compare the EVM bytecode language to the best-known bytecode: Java (JVM) bytecode—a much higher-level IR. The differences include:

- Unlike JVM bytecode, EVM does not have the notion of structs or objects, nor does it have a concept of methods.
- Java bytecode has a rich type system, EVM bytecode has a single type: a 256bit word.
- In JVM bytecode, stack depth is fixed under different control-flow paths: execution cannot reach the same program point with different stack sizes. In the EVM bytecode, no such execution constraints exist, which make the identification of standard control-flow constructs very hard.
- All control-flow edges (i.e., jumps) are to variables, not constants. The destination of a jump is a value that is read from the stack. Therefore, a value-flow analysis is necessary even to determine the connectivity of basic blocks. In contrast, JVM bytecode has a clearly-defined set of targets for every jump, independent of value flow.

- JVM bytecode has defined method invocation and return instructions. In EVM bytecode, although calls to outside a smart contract can be resolved, function calls inside a contract are translated to just jumps (to variable destinations, per the above point). All functions of a contract are fused in one stream of instructions, with low-level jumps as the means to transfer control.

To call an intra-contract function, the code pushes a return address to the stack, pushes arguments, pushes the destination block's identifier (a hash), and performs a jump (which pops the top stack element, to use it as a jump destination). To return, the code pops the caller's basic block identifier from the stack and jumps to it.

B. Declarative Program Analysis

Our work is based on declarative static program analysis, applied to smart contract decompilation. Declarativeness refers to implementing an analysis as a collection of logical rules (i.e., simple implications) that lead to inferences, which in turn trigger more rule inferences, up to the least fixpoint. The Datalog language is the most standard vehicle for declarative analysis approaches, both low-level [4], [16], [19], [27], [29], [36] and high-level [7], [11], [23], [24]. Additionally, Datalog-style inference rules have been used in formal specification tasks, such as the official specification of the Java VM verifier [17, p.170-320].

Datalog is less of a programming language and more of a specification language, since computation is based on two constructs: logical implication rules, and recursion. A Datalog rule “ $C(z,x) \leftarrow A(x,y), B(y,z).$ ” means that if $A(x,y)$ and $B(y,z)$ are both true, for some values x,y,z , then $C(z,x)$ can be inferred. Syntactically, the left arrow symbol (\leftarrow) is a logical implication symbol and separates the inferred facts (i.e., the *head* of the rule) from the previously established facts (i.e., the *body* of the rule).

Recursion is the standard vehicle for expressing computations in Datalog. Static analysis tasks are typically recursive, with multiple sub-algorithms collaborating towards a joint goal expressing the semantics of a program. The sub-algorithms have no clear stepwise order, but instead, all refer to (yet-incomplete) results of other sub-algorithm in a large recursive definition. In this way, the sub-algorithms enhance each other's results, each benefiting the others.

The declarative nature of the Datalog language means that *any* order of firing of the rules, and *any* order of evaluation of a rule's body will yield the same final result. To maintain this property, the implication in Datalog has to be strictly monotonic: each rule's firing can only introduce new facts and not hinder any of the previous inferences.

Part of the appeal of the Datalog language is that it currently enjoys several high-performance implementations including Soufflé [13]. Besides high-performance computation, state-of-the-art Datalog engines offer extensions to the language expressiveness: one can relax pure declarativeness and introduce ordering, as well as invent new values via *constructor* functions. We will refer to such non-purely-declarative facilities explicitly in our technical presentation.

III. GIGAHORSE DECOMPILE SPECIFICATION

In this section, we present the core building blocks of the Gigahorse decompiler. We use Datalog as a specification language for the decompiler. The Datalog rules are simplified but keep the essential complexity of their counterparts in the actual Gigahorse implementation. We attempt to provide as much technical detail as possible without sacrificing the readability and understandability of this paper. In this effort, the exposition will be uneven: we give in-depth details in the first few sub-sections, and elide technical details in later parts (of both Sections III and IV) when the reader will be able to fill the gaps.

A. Overview of Decompile Steps

We next summarize the main decompilation actions performed by the Gigahorse decompiler. For the sake of exposition, we describe a conceptual stepwise process, even though the decompiler specification is declarative and does not have an explicit order of operations. Starting from the original bytecode, the Gigahorse decompiler:

- 1) Finds basic block boundaries. The output to the next step is the *original bytecode*, split by basic blocks.
- 2) Performs local analysis of stack effects of basic blocks. The input of the next step attaches to the bytecode relevant summaries of stack effects per block.
- 3) Performs whole-contract context- and flow-sensitive dataflow analysis with on-the-fly *control-flow graph* (CFG) construction. This information is used to produce a 3-address IR, with global registers. We refer to this IR as *global 3-address IR*.
- 4) Infers function boundaries heuristically (i.e., entry and exit blocks, together with function calls) for public and private functions. The function boundaries enable the decompiler to transform the global CFG into local CFGs and a call graph.
- 5) Infers function arguments and return arguments for all functions, introduces fresh variables for these and performs an intra-procedural flow-sensitive analysis to infer the flow of these fresh variables. The output form is a *functional 3-address IR*, i.e., all variables are local variables and are scoped. Data is passed around functions through formal arguments, return arguments, or external constructs like storage and memory.

All the above steps work in concert to derive a high-level representation from the original bytecode. In this section, we focus on steps 2 and 3, and assume that the input is already parsed as statements in basic blocks (step 1). Section IV focuses on steps 4 and 5.

The original bytecode and our two IRs (the global 3-address IR and the functional 3-address IR) have common elements, but also differ in important ways. Throughout the descriptions of these, we will override certain concepts, such as statements or variables. For instance, variables in the global 3-address IR are global variables, whereas for the functional 3-address IR they are all local. The distinction between these should be clear from context.

S is a set of statement identifiers, $S \subseteq \mathbb{Z}$	
C is a set of constants, $C \subseteq \mathbb{Z}$	
V is a set of new variables	
B is a set of basic block identifiers	
I is a set of stack indices, from 0 to 1023	
PUSH($s : S, val : C$)	BINOP($s : S$)
DUP($s : S$)	POP($s : S$)
JUMPDEST($s : S$)	JUMPI($s : S$)
NEXT($prev : S, next : S$)	
BLOCK($stmt : S, block : B$)	
BLOCKHEAD($block : B, stmt : S$)	
BLOCKTAIL($block : B, stmt : S$)	
PUSHESANDPOPS($sOrB : S \cup B, pushes : \mathbb{N}, pops : \mathbb{N}$)	
LOCALDEFINES($stmt : S, var : V$)	
LOCALSTACKOUT($stmt : S, index : I, vOrI : V \cup I$)	
LOCALSTACKIN($stmt : S, n : I, vOrI : V \cup I$)	
VARIABLEVALUE($var : V, val : C$)	
NEWFRESHVAR ($stmt : S, n : \mathbb{N}$) = $variable : V$	

Fig. 1. Our domain, input relations representing the original bytecode, computed relations and outputs. The input relations encode program instructions and other environment information.

B. Input Language

Figure 1 describes the schema of the input and main intermediate relations, together with the domains of the program representation at this level of abstraction. Stack indices I by definition are between 0 and 1023, and we assume that all arithmetic operations on I are only defined in that range (i.e., no overflow or underflows). Notice that the original bytecode relations, such as PUSH or BINOP, make no references to variables since the EVM is a stack-based machine. Relations that capture instructions refer to a unique statement identifier. The PUSH instruction pushes a constant to the top of the stack (which starts from 0). The BINOP operation denotes any binary operation that takes the first two operands from the stack and returns a result into the stack. The specific form of binary operations is elided for presentation reasons, but examples include ADD, MUL, SHA, etc.

The JUMPDEST operation is a no-op, which only serves to mark its statement identifier as being a valid address to jump to. The JUMPI instruction is a conditional jump, which jumps to the address at the topmost element of the stack, if the element of stack position 1 is not zero. The input relation NEXT returns the next statement identifier in program order for a given statement identifier. The relation BLOCK maps a statement to its block, whereas BLOCKHEAD and BLOCKTAIL indicate the head and tail statements of a block. Finally, input relation PUSHESANDPOPS returns the maximum number of stack elements that a statement or an entire block pushes and pops. More precisely, for basic blocks this is a high-/low-watermark computation: the maximum and negated minimum, over all points in a block, of the balance of elements pushed minus those popped at that program point.

C. Local Stack Analysis

The bottom part of Figure 1 shows the relations inferred by step 2 of the decompilation: the local stack analysis. This

step summarizes the effects on the stack, per basic block, introduces the concept of variables (since none exist in the input representation), as well as performs a first value analysis, computing a static abstraction of the contents of the execution stack at every statement.

Relation **LOCALDEFINES** connects a statement with a freshly introduced variable, if the statement pushes new elements on the stack. Relations **LOCALSTACKIN** and **LOCALSTACKOUT** model the variables (or stack aliases) that each stack position contains before and after executing each statement, respectively. For instance, **LOCALSTACKOUT**(*stmt*, *index*, *v*) means that at statement *stmt*, stack position *index* contains the same value as variable *v* (freshly-introduced by some previous statement). The domain of these relations includes both variables and stack aliases ($V \cup I$). A stack alias (i.e., a stack index outside the range of values pushed by the current basic block) refers to a value that existed in a given stack position before the beginning of the basic block. This is because each basic block is analyzed in isolation, and stack values can also be passed from block to block. The earlier-defined input relation **PUSHESANDPOPS** gives the maximum number of possible stack aliases.

Finally, relation **VARIABLEVALUE** maps some of the variables to program constants. Any variable that is not present in this relation is considered dynamically computed, and its value is not modeled in this formalization (although we partially model some dynamic operations in the full implementation).

The rules that describe the local stack analysis are shown in Figure 2. The computation of **LOCALDEFINES** introduces new variables (one per statement that needs it) via a constructor **NEWFRESHVAR**, also used in later sections. This constructor function will typically be defined to retain all information passed to it:

NEWFRESHVAR(*stmt*, *n*) = *pair*(*stmt*, *n*) : *V*

The computation of relations **LOCALSTACKIN** and **LOCALSTACKOUT** is more involved, and the relations are mutually recursive. (As a convention, the rules use *** as a wildcard, i.e., it denotes any element, which is ignored.) Both relations have easy base cases: for **LOCALSTACKOUT**, if a statement assigns a fresh variable, then stack position 0 holds the value of this variable. For **LOCALSTACKIN**, if a statement is the first in a basic block, all stack positions that may be accessed contain stack aliases (i.e., their contents are represented by just the appropriate stack index and not by a fresh variable).

The recursive rules for **LOCALSTACKOUT** appeal to **LOCALSTACKIN**: At a **DUP**, the top of the stack after the instruction has the same value as the top of the stack before it. At all instructions, the unchanged contents of the stack are propagated, at indexes adjusted based on the pushes and pops performed by the instruction. Similarly, the recursive rule for **LOCALSTACKIN** merely propagates all values from the previous statement in the same basic block.

The above formulation illustrates well the declarative approach in the Gigahorse decompiler. Inspecting the rules and their apparent simplicity, it is hard to even convince oneself that they compute something useful. Yet they compactly capture the full flow of constant values over the EVM execution stack, at every stack position. Furthermore, the formulation

NEWFRESHVAR(*stmt*, 0) = *v*,
LOCALDEFINES(*stmt*, *v*) \leftarrow
PUSHESANDPOPS(*stmt*, *n*, ***), *n* > 0, !**DUP**(*stmt*).

LOCALSTACKOUT(*stmt*, 0, *var*) \leftarrow
LOCALDEFINES(*stmt*, *var*).

LOCALSTACKOUT(*stmt*, 0, *vOrI*) \leftarrow
DUP(*stmt*), **LOCALSTACKIN**(*stmt*, 0, *vOrI*).

LOCALSTACKOUT(*stmt*, *n* + *pushes* - *pops*, *vOrI*) \leftarrow
LOCALSTACKIN(*stmt*, *n*, *vOrI*),
PUSHESANDPOPS(*stmt*, *pushes*, *pops*), *n* \geq *pops*.

LOCALSTACKIN(*stmt*, *i*, *i*) \leftarrow
BLOCKHEAD(*block*, *stmt*),
PUSHESANDPOPS(*block*, ***, *pops*), *i* < *pops*.

LOCALSTACKIN(*stmt*, *n*, *var*) \leftarrow
LOCALSTACKOUT(*prevStmt*, *n*, *var*),
NEXT(*prevStmt*, *stmt*), !**BLOCKHEAD**(***, *stmt*).

VARIABLEVALUE(*var*, *val*) \leftarrow
PUSH(*stmt*, *val*), **LOCALDEFINES**(*stmt*, *var*).

Fig. 2. Relations to perform local (within each basic block) analysis

carefully mixes the concept of freshly-introduced variables with that of values at a stack position. The fresh variables both are essential for decompiling into a 3-address representation and serve as names to succinctly represent values.

D. Global Stack Analysis and CFG construction

Armed with a model of the contents for stack positions, the decompiler can now compute the targets of jump instructions (since these are values obtained from the stack). Global stack analysis and control-flow graph (CFG) construction is one of the most computationally-heavy parts of the analysis. We introduce additional relations in Figure 3 (top) and their definition (bottom).

Relations **BLOCKIN** and **BLOCKOUT** are global analogues of **LOCALSTACKIN** and **LOCALSTACKOUT**, and are computed by transferring values between each other according to the CFG defined by **GLOBALCFG**. Computing **GLOBALCFG** is the main part of this analysis step. A trivial CFG edge exists (in inverse direction) between the first statement of a basic block (given by **BLOCKHEAD**) and any of its predecessor statements. More interestingly, a CFG edge exists between any conditional jump instruction (**JUMPI**) and any basic block address (i.e., a constant value) held by the variable used to denote the jump target in the program. This makes CFG construction mutually recursive with the global flow analysis, as is typical with state-of-the-art frameworks for program analysis of higher-order languages with virtual calls, such as Java [28]. The final outputs of this analysis are embedded in relations **GLOBALCFG** and **LOCALUSES**. The latter relation represents which variable flows to which statement, in the same order as in the original EVM stack. The relation effectively connects the value flow between basic blocks, by resolving stack aliases

```

BLOCKIN(block : B, index : I, var : V)
BLOCKOUT(block : B, index : I, var : V)
JUMPTARGET(from : B, var : V, to : B)
GLOBALCFG(from : B, to : B)
LOCALUSES(stmt : S, position : N, var : V)


---


BLOCKIN(jumpee, index, var) ←
  BLOCKOUT(jumper, index, var),
  GLOBALCFG(jumper, jumpee).

BLOCKOUT(jumpee, index+pushes-pops, var) ←
  BLOCKIN(jumpee, index, var),
  PUSHESANDPOPS(jumpee, pushes, pops), index ≥ pops.

BLOCKOUT(block, index, var) ←
  GLOBALCFG(*, block), BLOCKTAIL(block, stmt),
  LOCALSTACKOUT(stmt, index, var), var ∈ V.

GLOBALCFG(jumper, ftBlock) ←
  GLOBALCFG(*, jumper), BLOCK(stmt, jumper),
  NEXT(stmt, ft), BLOCKHEAD(ftBlock, ft).

JUMPTARGET(jumper, var, jumpee),
GLOBALCFG(jumper, jumpee) ←
  BLOCK(stmt, jumper), JUMPI(stmt),
  LOCALUSES(stmt, 0, var),
  VARIABLEVALUE(var, jumpdest),
  JUMPDEST(jumpdest), BLOCK(jumpdest, jumpee).

LOCALUSES(s, i, v) ←
  LOCALSTACKIN(s, i, v), v ∈ V.

LOCALUSES(s, i, v) ←
  LOCALSTACKIN(s, i, j), j ∈ I,
  BLOCK(s, b), BLOCKIN(b, j, v).

```

Fig. 3. Rules to describe the global stack analysis and on-the-fly control-flow graph construction.

(i.e., stack positions set by predecessor basic blocks, as defined in Section III-C) when the interconnectivity of basic blocks is determined. This again demonstrates well the compactness and power of a declarative specification.

After performing these analyses, we can now produce global 3-address code using the schema and rules listed in Figure 4. This representation is adopted from the Vandal [9], [35] decompiler. The conversion to 3-address IR at this point is straightforward. Syntax sugar and minor detail elision are employed for presentation purposes. Language syntax is quoted using [and] and implicitly unquoted for meta-variables. For instance, $s:[to := \text{BINOP}(x, y)]$ indicates that statement s is some binary operation on x and y with its result in to , where x , y , and to are the meta-variables referring to the bytecode variables. The distinction between variables in the analyzed program and meta-variables in the analysis is clear from context, therefore we simply refer to “variables”. We reuse the instruction opcodes from the original bytecode whenever these make sense. Instructions that do not define any new variables or perform computation, such as DUP or JUMPDEST are not present in this representation. From this point, unless otherwise

```

s:[to := CONST(c)] ←
  PUSH(s, c), LOCALDEFINES(s, to).

s:[JUMPI(cond, label)] ←
  JUMPI(s),
  LOCALUSES(s, 0, cond), LOCALUSES(s, 1, label).

s:[to := BINOP(a, b)] ←
  BINOP(s), LOCALDEFINES(s, to),
  LOCALUSES(s, 0, a), LOCALUSES(s, 1, b).

```

Fig. 4. The rules and relations that compute the global 3-address IR, an important intermediate decompiler representation.

specified, whenever we refer to “statements”, or their identifier S , we will be referring to statements in the 3-address IR representation.

IV. RECONSTRUCTING SOURCE LEVEL FUNCTIONS

An essential element in the Gigahorse decompilation is the inference of functions, which have been dissolved by the compilation to EVM bytecode. We find that this inference process enters a virtuous circle with the CFG inference of the previous section: CFG inference informs function reconstruction, and functions inform an even more precise CFG inference.

A. Public Functions

The current EVM version does not have notions of functions (call stacks will be introduced in a future version [30]). However, all programming languages that compile down to EVM support functions. The ABI (application binary interface) for Solidity and the EVM introduces the notion of public functions, so that external contracts or users can call public functions by specifying the hash of the function signature. By identifying the dispatching patterns that are introduced by the Solidity compiler, the Gigahorse decompiler finds all public function entries in the code. For instance, let us look at the following simplified global 3-address IR:

```

// get hash supplied from caller
v12 = CALLDATALOAD v10
// check whether supplied hash matches
v1a = EQ v12 0x6fdde03
// conditionally jump to function 0x6fdde03
JUMPI 0x139 v1a
// nope, try this one instead
v25 = EQ v12 0x95ea7b3
// conditionally jump to function 0x95ea7b3
JUMPI 0x1b4 v25

```

We can see that there are two public functions, one with the function descriptor of 0x6fdde03 that begins at address 0x139, and one with the function descriptor 0x95ea7b3 that begins at address 0x1b4. Since the ABI is standard, we expect that all dispatching patterns by all compilers are catered for. This matching is rather straightforward for any decompiler to perform.

The only additional feature that Gigahorse adds is a presentational convenience: Gigahorse tries to match function descriptors, such as 0x95ea7b3, to the source-level function signatures, such as `approve(address, uint256)`. This is made possible by consulting an existing database [31]

of public function signatures. This allows detecting source signatures for over two-thirds of public functions in deployed contracts.

B. Searching for Functions that Return

The challenge in reconstructing functions concerns private functions, which have all but disappeared from the compiled code. To detect private functions, Gigahorse employs complex heuristics that require a full global dataflow information propagation and the construction of CFGs. The first heuristic is to look for specific instances of passing addresses on the stack inter-procedurally, when these addresses are subsequently used for further jumps after running a function. For instance, let us look at the following simplified global 3-address IR:

```
bar: ...
    va = CONST <ret> // set return address
    vb = CONST 0xFF // set data
    vc = CONST <foo> // set function address
    JUMPI vc 0x1 // jump to foo
ret: ...
foo: ...
    JUMPI va 0x1 // jump to 'ret'
```

This program snippet first passes the return address `<ret>` before jumping to `<foo>`. At the end of `<foo>`, it jumps back to the return address that was passed before. This function search heuristic, therefore has to identify that (a) a basic block (return) jumps to a valid non-locally-derived address, which (b) originates at another block (the caller) that can reach the return basic block. This heuristic is condensed in a relation, which is further refined into `FNCALLRET` by:

- Making sure that a basic block can only belong in a single function. When a conflict arises, we deterministically pick a function according to a total ordering of our choice.
- Making sure that a function can only be entered through a function call.

C. Decomposing Functions Further

Detecting functions that return to their callers is not sufficient in the context of the EVM. Functions that terminate (e.g., `halt`) are disproportionately common in smart contracts, since smart contract computation is typically short. Therefore, Gigahorse employs additional logic to detect functions that do not return. The main idea is that a basic block is in an independent function if it is reachable from two (or more) previously-identified distinct functions. This is an intuitively inevitable rule: if two functions both use the same code, this code must also be factored into a reusable fragment, i.e., a function. (Note that, unlike the technique of Section IV-B, this approach only detects source-level functions that are called more than once—functions that are called just once are inlined with no loss of precision.)

In more detail, the detection logic identifies basic blocks reachable via paths that: (a) start from two or more function entries (corresponding to previously-identified functions A and B, for instance), where (b) each path remains within its source function (A or B respectively). The process is actually recursive—as more functions are discovered, more opportunities arise for some basic block to be reachable from

more functions. Although the number of functions discovered is monotonically growing, the logic is not monotonic at the Datalog syntax level: in order to describe paths that occur strictly within a single function, the logic needs to express that “a path does not cross into a different function”, i.e., to use the negation of the call-graph edges predicate, whose contents are also growing in the same computation.¹ For this reason, we use a fixpoint loop external to the (monotonic) Datalog rules, so that we are able to iteratively recompute relations at every step, and each relation can refer to versions at the previous iteration. (In the implementation, this is done via the standard “components” facility of the Soufflé Datalog engine [13] that Gigahorse uses.)

The complete algorithm is shown in Figure 5. The input relation, `PUBLICFUNCTIONCALL` is derived as described in Section IV-A, while `FNCALLRET` is described in Section IV-B. From these two relations we compute the inputs our algorithm, i.e., `CALLGRAPH0` and `FUNCTIONENTRY0`. The algorithm then proceeds to discover new functions at each iteration (`FUNCTIONENTRYn` and `CALLGRAPHn`), and each time recomputes paths reachable from the function entry (`REACHFROMn`).

After computing a least fixpoint on `FUNCTIONENTRY`, i.e., no more new functions are discovered, we can proceed to compute local CFGs and call graphs. The call graphs are computed by taking a union over all `CALLGRAPHn`, i.e., $\bigcup_{n \in \mathbb{N}} \text{CALLGRAPH}_n$, which we simply refer to as `CALLGRAPH` from this point onwards. We also assume the same for `FUNCTIONENTRY`. The local CFGs are then computed as follows:

$$\frac{\text{LOCALCFG}(\text{block} : B, \text{next} : B)}{\text{LOCALCFG}(\text{block}, \text{next}) \leftarrow \text{GLOBALCFG}(\text{block}, \text{next}), \text{!FUNCTIONENTRY}(\text{next}), \text{!FNCALLRET}(*, *, \text{block}, \text{next}).}$$

$$\text{LOCALCFG}(\text{block}, \text{next}) \leftarrow \text{GLOBALCFG}(\text{block}, \text{func}), \text{FNCALLRET}(\text{block}, \text{func}, *, \text{next}).$$

That is, the function inference helps *filter out* previously-inferred CFG edges, if these do not agree with the functional abstraction. Such spurious edges arise due to inherent imprecision in any static analysis. In our rules, this imprecision is mainly introduced at control-flow join points, i.e., because of multiple predecessors of a basic block getting their stack contents mixed up, based on the logic of Figure 3.

D. Inferring Function Arguments and Local Variables

Inferring function boundaries and local CFGs is very beneficial to client analyses. Not only are local CFGs more precise (as discussed) but function-level reasoning enables summary-based analyses (which are typically very scalable and highly precise). In order to enable summary-based data-flow analyses, we need to have local variables, and minimize the number of global variables.

¹No real non-monotonicity exists: even though some path may later be found to be invalid—by crossing into a different function—different paths with the same overall property inevitably exist.

```

PUBLICFUNCTIONCALL(caller : B, func : B)
FNCALLRET(caller : B, func : B, ret : B, rTarg : B)
RETURNBLOCK(ret : B)
CALLGRAPHn:N(caller : B, func : B)
FUNCTIONENTRYn:N(func : B)
REACHFROMn:N(func : B, block: B)
RETURNBLOCK(ret),
FUNCTIONENTRY0(func),
CALLGRAPH0(caller, func) ←
  FNCALLRET(caller, func, ret, *).

FUNCTIONENTRY0(func),
CALLGRAPH0(caller, func) ←
  PUBLICFUNCTIONCALL(caller, func).

FUNCTIONENTRY0(0x0).
i = 1.
DO {
  REACHFROMi(block, block) ←
    FUNCTIONENTRYi-1(block)

  REACHFROMi(func, next) ←
    REACHFROMi(func, block),
    GLOBALCFG(block, next),
    !CALLGRAPHi-1(block, next),
    !RETURNBLOCK(block).

  CALLGRAPHi(prev, block),
  FUNCTIONENTRYi(block) ←
    REACHFROMi(f1, block), REACHFROMi(f2, block),
    GLOBALCFG(prev, block), !RETURNBLOCK(prev),
    !REACHFROMi(f1, prev), !REACHFROMi(f2, prev).

  i = i + 1.
} UNTIL FIXPOINT(FUNCTIONENTRY)

```

Fig. 5. Heuristic for decomposing functions further.

Gigahorse infers the *number* of function arguments by computing the number of caller-supplied elements that the entire function pops from the stack throughout its execution. Similarly, the number of function return arguments is inferred by calculating the balance of extra (pushed and not popped) elements pushed up to a call instruction. Unfortunately, the EVM does not guarantee that the stack depth is statically known at each program point (unlike, say, the Java VM). A good practical solution to this is to infer all possible push and pop balances along a function’s execution paths, up to a finite upper bound, and take the minimum number of these at the end of the returning basic block. This is a heuristic strategy, whose effectiveness is validated experimentally.

After inferring the number of function arguments and return values, Gigahorse proceeds to create fresh variables for them using the **NEWFRESHVAR** constructor. This logic is elided for space reasons.

The last analysis step is to infer (non-argument) local variables and places where these variables flow to. This is a computation analogous to the introduction of variables, before functions are detected, in Section III-D (predicate

LOCALUSES in Figure 4). The main output of the analysis is predicate FUNCTIONALUSES($s : S, i : I, v : V$), which, for each statement, indicates which local variables are used and in which order. Using this relation, we can produce the final functional 3-address IR, similar to the process described in Figure 4. At this point, we can also introduce the additional opcodes PRIVATECALL and PRIVATERETURN, which substitute jump instructions that call functions or return from a function.

V. IMPLEMENTATION AND DISCUSSION

The specification of the previous sections captures the essence of the Gigahorse implementation. The actual Datalog specification of the decompiler has more technical details, comprising several hundred logical rules (over 3K lines of Datalog, using Soufflé [13] as the dialect and execution engine) and a small Python scaffolding (of around 1K lines) borrowed from Vandal [35]. Compared to the specification, the full implementation contains:

- handling of the full instruction set of the EVM, as opposed to the minimal instruction set presented;
- several more decompilation heuristics, secondary but complementary to the ones discussed in Sections III and IV.
- context sensitivity throughout the rules: concepts such as CALLGRAPH have their elements qualified by a “context”, which allows more precise static analysis. (We use a 1-call-site context—a.k.a. 1-CFA—as the default.)
- a library facilitating further client analyses;
- example clients.

We next discuss some of the above in more detail.

A. Deployment

Gigahorse has been applied over the entire contents of the Ethereum blockchain, with the results of decompilation offered as the free Contract Library service (<https://contract-library.com>). Contract Library has received very positive feedback from the Ethereum community.

B. Output and Client Analyses

Gigahorse produces output both in structured form (i.e., tables, exported as CSV files) and in pretty-printed text. A very simple contract, below, helps as illustration.

```

function _functionSelector() public {
  v27 = (CallDataLoad(0x0) / 0x1000...);
  if (0x5fdf05d7 == v27) two();
  if (0x901717d1 == v27) one();
  exit();
}
function two() public { f0x57(0x2); exit(); }
function one() public { f0x57(0x1); exit(); }
function f0x57(varg1) private {
  STORAGE[0x0] = varg1; return();
}

```

Gigahorse infers four functions in total, two of which are public functions with a high-level function name. A private function, `f0x57`, is also inferred. All low-level jumps have been transformed into high-level control-flow and calls/returns. Of the jump instructions in the original contract bytecode, only the one corresponding to the return statement is polymorphic,

since it can return to two callers. Gigahorse has no trouble detecting this as a proper function with return, which enables more precise data flow. The other decompilers in our evaluation set, Porosity, Vandal and EthIR, fail to find the private function and call-return pattern.

Gigahorse has been designed as a framework for writing security-related analyses on top of its high-level functional 3-address IR. In order to facilitate the development of these client analyses, Gigahorse offers additional development tools. For client analyses implemented in Datalog, we developed a bulk analyzer that can be supplied with a user-defined pipeline of client analyses. These are compiled and executed in parallel, in tandem with the decompilation. Client analyses for Gigahorse can be written in any language by reading the decompilation results from CSV files. In fact, one of our client analyses is a high-level pretty-printer, which takes the decompiler’s output and outputs high-level Solidity-like code, including function signatures, control-flow structures such as `if` statements, some types and complex expressions. This is a useful tool for the human inspection of smart contracts.

Additionally, Gigahorse offers a Datalog API for the decompilation results together with libraries of analysis functions tuned for the decompiler’s output. These libraries perform data-flow, dependency, and loop-semantic analysis. The data-flow analysis library is able to compute an intra-procedural data-flow analysis or summary-based inter-procedural data-flow analysis. The latter is enabled by the rich functional decompilation described in the previous section. All analyses are highly parametric. Users can instantiate multiple versions, each defining its own transfer functions, and can also limit the scope of an analysis (e.g., analyze only loop-induction variables within loops). The loop-semantic analysis identifies loops and other control-flow structures, their exit conditions, induction variables, dominators, etc. Using the provided libraries we have ported the recently-published MadMax client analysis for gas-related vulnerabilities [9] to the Gigahorse decompiler with similar results but with much higher performance, and for over 99.9% of the deployed smart contracts. The re-implementation of MadMax comprises just 250 lines of code, showcasing the power of the client analysis infrastructure of Gigahorse.

C. Declarativeness and Monotonicity

The declarative nature of Gigahorse has been a significant facilitator of its decompilation effectiveness. (Speculatively, this is an insight that may be also applicable in other decompilation domains, such as machine code decompilation.) Logical rules allow for concise expression of decompilation patterns and heuristics. Perhaps more importantly, separate patterns can be specified completely independently, with no ordering or other artificial dependency. Still, the independent patterns can benefit or complement each other. Dependencies between decompilation patterns (e.g., that one needs to run before another) are determined automatically by the execution engine. Dependencies between decompilation patterns arise naturally in the specification we saw in earlier sections—e.g., the call-graph informs the flow analysis and vice versa. We also saw a representative example of complementary decompilation

patterns in Section IV: one pattern detects functions by tracking that they are called and return to their call site), while another detects that a block must be in a separate function because it is reachable by two existing ones. Neither pattern is more complete or more precise than the other.

VI. EVALUATION

There are several research questions that our experiments intend to answer:

RQ1: Scalability—is Gigahorse a scalable and efficient decompiler, for all contracts in the wild?

RQ2: Completeness/Coverage—how well does the decompiled code cover the original code available in the wild?

RQ3: Precision—does the decompiled code precisely match high-level semantics?

For most of these research questions, we will be comparing Gigahorse against Porosity (the best-known, state-of-the-practice Ethereum decompiler at submission time) and Vandal [35] (a recent research tool, probably the closest comparable in the literature). In Section VI-E we also perform an experiment with less closely related tools.

Our experimental setup consists of all programs available on the Ethereum blockchain as of April 2018. This makes up the universe set of “contracts in the wild” of around 6.6 million contracts deployed from 91.8K unique programs.

We ran all systems on an idle machine with an Intel Xeon E5-2687W v4 3.00GHz and 512GB of RAM. To bound the (long) time to run experiments, we set a cutoff of 120s for decompilation. (As we show, decompilation time for successfully-decompiled contracts is on average over 10 times below this threshold, for all systems.) Any contracts that take longer to decompile are considered to timeout. All experiments were conducted in parallel. Our machine has 48 logical cores, so we ran 45 processes in parallel. Note that the Gigahorse decompiler can also be configured to parallelize decompilation within a *single* contract, which is an option we did not enable when analyzing multiple contracts in bulk.

We emphasize that our evaluation is quantitative, based on several metrics, but we have confirmed by manual inspection of numerous contracts that the metrics reflect well the actual quality of decompilation.

A. RQ1: Scalability

To answer this question we employ the following metrics:

Timeouts/fatal errors, i.e., proportion of contracts that were not successfully decompiled due to timeouts or fatal errors.

Contracts with at least one function detected according to the respective decompiler.

Time i.e., wall-clock time taken to decompile the average contract, excluding contracts that time out.

Contract size in terms of the number of functions computed by the Gigahorse decompiler.

The table below summarizes the first three metrics. (We use the convention that for metrics in *italics* lower is better, otherwise higher is better.)

	no timeout	≥ 1 functions	<i>time</i>
Gigahorse	99.98%	79.9%	0.7s
Vandal	88.13%	63.7%	5.7s
Porosity	48.56%	17.4%	1.2s

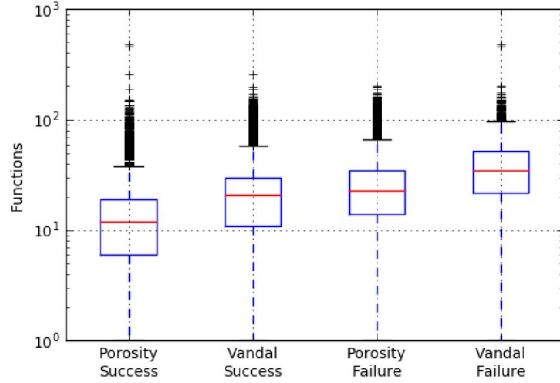


Fig. 6. Successful decompilation vs. unsuccessful decompilation against size of contracts (number of functions).

Gigahorse is both very fast (over 4x faster than Vandal, for higher decompilation quality) and scalable to all contracts. Both Porosity and Vandal fail to decompile a significant portion of the contracts. As can be seen, Porosity is bimodal: it is fast for the contracts it manages to decompile, or hits a scalability wall. Our setup, giving Porosity 100x the time of its average successful decompilation (120s), still did not allow decompilation of more than half of the contracts. The contracts that are not decompiled correctly are larger in size, for both Porosity and Vandal—Figure 6 plots decompilation success vs. contract size in functions (as reported by Gigahorse, whose function inference is the most reliable).

Gigahorse also has the highest proportion of contracts with at least a single function detected (apart from the dispatcher function). Having 0 functions is not necessarily proof of incomplete decompilation, but it is a strong indicator, especially when the numbers for Porosity and Vandal diverge from Gigahorse, whose function inference is very reliable.

B. RQ2: Completeness

We employ the following metrics:

Unknown jump targets, i.e., proportion of jumps that do not have a target for the non-fallthrough case.

Unreachable code, i.e., % of unreachable basic blocks.

Number of functions per contract, as indicated by the respective decompiler.

The results of these metrics are shown in the following table.

	unknown jump targets	unreachable code	functions / contract
Gigahorse	0.00%	8.7%	21.0
Vandal	0.81%	19.8%	12.2
Porosity	N/A	N/A	1.92

Note that the output of Porosity is not in a format that enables the measurement of jump targets of unreachable code, so we only measure functions. Gigahorse finds jump targets for virtually all jump instructions, which translates into less unreachable code than Vandal. The number of functions detected (public or private) is also significantly higher.

C. RQ3: Precision

To answer the question we employ the following metrics:

Polymorphic jumps, i.e., proportion of jumps that resolve to more than one target for non-fallthrough CFG edges.

Unstructured loops, i.e., proportion of loops with unusual control-flow structure. This indicates imprecision in the control-flow graph.

Loops with an exit condition—the converse likely indicates imprecision in the control-flow graph.

Data flows (values) per variable, in arithmetic or assignments.

We only compare Gigahorse against Vandal in this case, as Porosity does not enable measurement of these metrics.

	Poly. jumps	Unstr. loops	Loops w/exit	Data flows
Gigahorse	2.58%	0.00%	96.6%	3.82
Vandal	2.72%	11.7%	91.7%	8.20

Gigahorse fares better in all metrics. For instance, it computes fewer polymorphic jumps (i.e., jumps with multiple possible targets), primarily due to better context sensitivity during the whole-contract analysis. The “data flows” and “unstructured loops” metrics are most telling, and we discuss them next in the context of client analyses.

D. Discussion: Benefits for client analyses

RQ3 shows that the number of values flowing to variables at arithmetic operations, is much lower for Gigahorse. This is due to a variety of factors: summary-based analysis, context sensitivity, and more precise CFGs. The data-flow relation is a client of the decompiler in its own right, but it is also useful for higher-level clients, such as vulnerability detection analyses.

The detection of loops is also a client analysis offered by the Gigahorse toolchain, where we can see that there are fewer unstructured loops detected in the inferred CFGs. Since there are no high-level constructs in Solidity to write unstructured loops, we expect that a good quality CFG has virtually no unstructured loops (which is the case in Gigahorse). Here, the context sensitivity and local CFGs offered by Gigahorse directly translate into a higher-quality client analysis.

We ported the recent MadMax gas-vulnerability detector [9] (the main client of the Vandal decompiler) to Gigahorse. Gigahorse analyzes virtually all contracts, rather than 92% in the case of Vandal. The performance of the client analysis also benefits—the additional MadMax client analysis only costs 0.15s of wall clock time per contract on average. The table below summarizes the vulnerabilities flagged by MadMax with decompilation under Vandal vs. under Gigahorse. (Neither higher nor lower numbers are better by default for this table.)

	Unbounded Iteration	Overflow Loop Iter.	Wallet Griefing
Gigahorse	3.9%	1.8%	0.32%
Vandal	4.1%	1.2%	0.12%

Although for liberal analyses (unbounded iteration) the improved precision of Gigahorse yields lower numbers, for more exotic vulnerability patterns the improved completeness of Gigahorse (managing to analyze some-10% more contracts, among the largest available) yields more warnings.

E. Sensitivity Analysis

We next discuss varying several dimensions of our experiments and how this affects the results.

a) Different Vandal settings: Published work on Vandal [9] claims a higher percentage of successfully decompiled contracts: 91.8% vs. 88.1% in our experiments. The numbers are close and the difference is due to slightly different settings. To confirm this, we also tried a different Vandal setup, for maximum decompilation ability, and got it to succeed in even more contracts—94.7%—but at the expense of significantly worse precision than that of our experiments.

b) Other Research Decompilers: In addition to the pure decompilers we compare against, several other recent research tools perform some form of decompilation before further analysis. We do not attempt a full comparison with such tools, but performed a more limited experiment of their end-to-end scalability and generality, in the same setup as for RQ1. We consider EthIR [1], Mythril [22], and Rattle [34]. (Of these, EthIR is the closest to a pure decompiler.) EthIR can handle 73% of the 91K unique contract programs, for an average time of 11.9s per contract. Most of the failing contracts produce an error-state exit. Our own inspection of the rule-based representation output confirms that EthIR only performs simple local reasoning within each basic block. Mythril can handle 40% of 91K unique contracts with an average time of 19.1s per contract; exiting with an error state for 19% of contracts. On the other hand Rattle can handle 69.7% of the 91K unique contracts with the same 120s timeout. It produced non-timeout error-state exits on 24% of the contracts.

In recent months, the Eveem system [15] has emerged as a decompiler for the Ethereum blockchain. Eveem offers a convenient query language and high responsiveness, for a positive user interaction experience. Eveem is based on a symbolic execution engine, much like other smart contract tools—e.g., EthIR. Compared to static analysis, a symbolic execution approach lags in completeness: it may miss significant parts of the code, due to not finding symbolic inputs that can exercise them. (At the same time, symbolic execution may yield higher precision for the code it does decompile.) A recent anecdote illustrates this lag. In Feb. 2019 electronic discussions (Security Community Telegram group, Feb.10), the question “how many deployed contracts call the 3 blockchain-compiled functions” was posed. Eveem produced an answer of 8,000 whereas Gigahorse returned over 40,000 contracts, showcasing its much higher code coverage.

VII. RELATED WORK

Analysis and verification for smart contracts has received substantial attention recently due to the security issues inherent in the high-risk paradigm of smart contract development.

a) Decompilers for smart contracts: Vandal [9], [35] is an open-source framework written in Python, and consists of a decompiler, a blockchain scraper, and a set of extensible vulnerability analyses written in Soufflé [13] Datalog. We have compared Gigahorse to Vandal extensively in previous sections. Porosity [33] is a high-level decompiler from EVM bytecode to Solidity-like source (similar to our pretty-printer output) implemented in C++. The EthIR [1] framework for high-level analysis of Ethereum bytecode based on the trace-based Oyente tool [18]. Its decompilation output introduces variables that are local to each basic block which

makes the analysis trivial. Note that the EthIR framework reconstructs some high-level control and data-flow fragments from Oyente traces. Fragments of the control-flow graph that are not covered by Oyente’s traces remain undiscovered by EthIR. Mythril is a security analysis tool for Ethereum smart contracts [22]. Mythril performs decompilation aided by a symbolic execution engine (Laser-EVM). It produces traces that are used to generate an intermediate representation and it therefore suffers similar incompleteness issues as EthIR. Similarly, Rattle [34] also constructs an IR in SSA form [6] and performs program analysis on it.

b) Analysis frameworks for smart contracts: Previous work on smart contract security analysis can be classified according to its underlying techniques, including symbolic execution, formal verification, and abstract interpretation. Some of the work does not necessarily need decompilation. Systems including Oyente [18], MAIAN [25], GASPER [5] and Grossman et al.’s recent work [10] use a symbolic execution/trace semantics approach. Thus, they analyze a single execution trace only. The formal verification tool by Bhargavan et al. [3] detects vulnerabilities that include not checking the return value of external address calls, and reentrancy using F*. Similarly, the FSolidM framework [20] checks for reentrancy and transaction ordering vulnerabilities. It can also detect coding patterns such as time-constraint and authorization issues as outlined in previous work [2]. The MAIAN framework [25] finds contract vulnerabilities including locking of funds indefinitely, leaking funds to arbitrary users, and killable smart contracts. The ZEUS system [14] conducts policy checking for a set of policies, including reentrancy, unchecked send, failed send, integer overflow, transaction state dependence/order, and block state dependence.

c) Decompilation for Java bytecode: There is a cornucopia of Java decompilers [8], [12], [21], [26], [32]. The Krakatoa decompiler [26] uses a pipeline of transformations to recover Java programs from Java bytecode. Gigahorse employs similar techniques to convert stack-based operations to three-address code via symbolic execution. In [8], Prolog is used to specify a Java bytecode decompiler. An evaluation paper [12] compared the performance of various Java decompilers.

VIII. CONCLUSION

We presented Gigahorse, a toolchain for decompiling and analyzing Ethereum smart contract binaries. Gigahorse leverages declarative program analysis to produce a highly-effective decompiler. The decompilation technology is unique due to the feedback loops between different abstraction levels. We have also shown that Datalog is very well suited for describing such mechanisms succinctly and without sacrificing performance. Furthermore, the Gigahorse toolchain is built as a full-fledged framework containing highly-parametric program analysis libraries and security client analyses.

ACKNOWLEDGMENTS

This research was supported partially by the Australian Government through the Australian Research Council’s Discovery Projects funding scheme (project ARC DP180104030).

REFERENCES

- [1] Elvira Albert, Pablo Gordillo, Benjamin Livshits, Albert Rubio, and Ilya Sergey. Ethir: A framework for high-level analysis of ethereum bytecode. In *Automated Technology for Verification and Analysis (ATVA)*. Springer, 2018.
- [2] Massimo Bartoletti, Salvatore Carta, Tiziana Cimoli, and Roberto Saia. Dissecting ponzi schemes on ethereum: identification, analysis, and impact, 2017.
- [3] Karthikeyan Bhargavan, Antoine Delignat-Lavaud, Cédric Fournet, Anitha Gollamudi, Georges Gonthier, Nadim Kobeissi, Natalia Kulatova, Aseem Rastogi, Thomas Sibut-Pinote, Nikhil Swamy, and Santiago Zanella-Béguelin. Formal verification of smart contracts: Short paper. In *Proceedings of the 2016 ACM Workshop on Programming Languages and Analysis for Security, PLAS '16*, pages 91–96, New York, NY, USA, 2016. ACM.
- [4] Martin Bravenboer and Yannis Smaragdakis. Strictly declarative specification of sophisticated points-to analyses. In *submission to OOPSLA '09: 24th annual ACM SIGPLAN conference on Object Oriented Programming, Systems, Languages, and Applications*, New York, NY, USA, 2009. ACM.
- [5] T. Chen, X. Li, X. Luo, and X. Zhang. Under-optimized smart contracts devour your money. In *2017 IEEE 24th International Conference on Software Analysis, Evolution and Reengineering (SANER)*, pages 442–446, Feb 2017.
- [6] Ron Cytron, Jeanne Ferrante, Barry K. Rosen, Mark N. Wegman, and F. Kenneth Zadeck. Efficiently computing static single assignment form and the control dependence graph. *ACM Trans. Program. Lang. Syst.*, 13(4):451–490, October 1991.
- [7] Michael Eichberg, Sven Kloppenburg, Karl Klose, and Mira Mezini. Defining and continuous checking of structural program dependencies. In *ICSE '08: Proceedings of the 30th International Conference on Software Engineering*, pages 391–400, New York, NY, USA, 2008. ACM.
- [8] Miguel Gómez-Zamalloa, Elvira Albert, and Germán Puebla. Decompilation of java bytecode to prolog by partial evaluation. *Inf. Softw. Technol.*, 51(10):1409–1427, October 2009.
- [9] Neville Grech, Michael Kong, Anton Jurisevic, Lexi Brent, Bernhard Scholz, and Yannis Smaragdakis. Madmax: Surviving out-of-gas conditions in ethereum smart contracts. *Proc. ACM Programming Languages*, 2(OOPSLA):to appear, November 2018.
- [10] Shelly Grossman, Ittai Abraham, Guy Golan-Gueta, Yan Michalevsky, Noam Rinetzk, Mooly Sagiv, and Yoni Zohar. Online detection of effectively callback free objects with applications to smart contracts. *Proc. ACM Programming Languages*, 2(POPL):48:1–48:28, December 2017.
- [11] Elnar Hajiye, Mathieu Verbaere, and Oege de Moor. Codequest: Scalable source code queries with Datalog. In *ECOOP'06: Proceedings of the 20th European Conference on Object-Oriented Programming*, pages 2–27. Springer, 2006.
- [12] James Hamilton and Sebastian Danicic. An evaluation of current java bytecode decompilers. In *Proceedings of the 2009 Ninth IEEE International Working Conference on Source Code Analysis and Manipulation, SCAM '09*, pages 129–136, Washington, DC, USA, 2009. IEEE Computer Society.
- [13] Herbert Jordan, Bernhard Scholz, and Pavle Subotić. Soufflé: On synthesis of program analyzers. In Swarat Chaudhuri and Azadeh Farzan, editors, *Computer Aided Verification*, pages 422–430, Cham, 2016. Springer International Publishing.
- [14] Sukrit Kalra, Seep Goel, Seep Goel, and Subodh Sharma. Zeus: Analyzing safety of smart contracts, 2018.
- [15] Tomasz Kolinko. Eveem/Panoramix – Showing Contract Sources since 2018, 2018. Accessed: 2019-02-15.
- [16] Monica S. Lam, John Whaley, V. Benjamin Livshits, Michael C. Martin, Dzintars Avots, Michael Carbin, and Christopher Unkel. Context-sensitive program analysis as database queries. In *Proc. of the 24th Symp. on Principles of Database Systems, PODS '05*, pages 1–12, New York, NY, USA, 2005. ACM.
- [17] Tim Lindholm, Frank Yellin, Gilad Bracha, and Alex Buckley. *The Java Virtual Machine Specification, Java SE 8 Edition*. Addison-Wesley Professional, 1st edition, 2014.
- [18] Loi Luu, Duc-Hiep Chu, Hrishi Olickel, Prateek Saxena, and Aquinas Hobor. Making smart contracts smarter. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security, CCS '16*, pages 254–269, New York, NY, USA, 2016. ACM.
- [19] Magnus Madsen, Benjamin Livshits, and Michael Fanning. Practical static analysis of JavaScript applications in the presence of frameworks and libraries. In *Proceedings of the ACM SIGSOFT International Symposium on the Foundations of Software Engineering*, August 2013.
- [20] Anastasia Mavridou and Aron Laszka. Designing secure ethereum smart contracts: A finite state machine based approach, 2018.
- [21] Jerome Miecznikowski and Laurie J. Hendren. Decompiling java bytecode: Problems, traps and pitfalls. In *Proceedings of the 11th International Conference on Compiler Construction, CC '02*, pages 111–127, London, UK, UK, 2002. Springer-Verlag.
- [22] Bernhard Mueller. Smashing ethereum smart contracts for fun and real profit, 2018. The 9th annual HITB Security Conference.
- [23] Mayur Naik, Alex Aiken, and John Whaley. Effective static race detection for java. In *Proc. of the 2006 ACM SIGPLAN Conf. on Programming Language Design and Implementation, PLDI '06*, pages 308–319, New York, NY, USA, 2006. ACM.
- [24] Mayur Naik, Chang-Seo Park, Koushik Sen, and David Gay. Effective static deadlock detection. In *Proc. of the 31st International Conf. on Software Engineering, ICSE '09*, pages 386–396, New York, NY, USA, 2009. ACM.
- [25] Ivica Nikolic, Aashish Kolluri, Ilya Sergey, Prateek Saxena, and Aquinas Hobor. Finding the greedy, prodigal, and suicidal contracts at scale. *CoRR*, abs/1802.06038, 2018.
- [26] Todd A. Proebsting and Scott A. Watterson. Krakatoa: Decompilation in Java (does bytecode reveal source?). In *Proceedings of the 3rd Conference on USENIX Conference on Object-Oriented Technologies (COOTS) - Volume 3, COOTS'97*, pages 14–14, Berkeley, CA, USA, 1997. USENIX Association.
- [27] Thomas Reps. Demand interprocedural program analysis using logic databases. In R. Ramakrishnan, editor, *Applications of Logic Databases*, pages 163–196. Kluwer Academic Publishers, 1994.
- [28] Yannis Smaragdakis and George Balatsouras. Pointer analysis. *Foundations and Trends in Programming Languages*, 2(1):1–69, 2015.
- [29] Yannis Smaragdakis, George Kastrinis, and George Balatsouras. In-trospective analysis: Context-sensitivity, across the board. In *Proc. of the 2014 ACM SIGPLAN Conf. on Programming Language Design and Implementation, PLDI '14*, pages 485–495, New York, NY, USA, 2014. ACM.
- [30] Various. Eip 214 - new opcode staticcall. <https://github.com/ethereum/EIPs/blob/master/EIPS/eip-214.md>. Accessed: 2018-08-15.
- [31] Various. Ethereum function signature database. <https://www.4byte.directory/>. Accessed: 2018-08-15.
- [32] Various. JODE – Java Optimize and Decompile Environment, 2018. Accessed: 2018-08-24.
- [33] Various. Porosity – a decompiler for EVM bytecode into readable Solidity-syntax contracts, 2018. Accessed: 2018-08-24.
- [34] Various. Rattle – An EVM Binary Static Analysis Framework, 2018. Accessed: 2018-08-24.
- [35] Various. Vandal – A Static Analysis Framework for Ethereum Bytecode, 2018. Accessed: 2018-07-30.
- [36] John Whaley, Dzintars Avots, Michael Carbin, and Monica S. Lam. Using Datalog with binary decision diagrams for program analysis. In *Proc. of the 3rd Asian Symposium on Programming Languages and Systems*, pages 97–118, 2005.